

System Programming

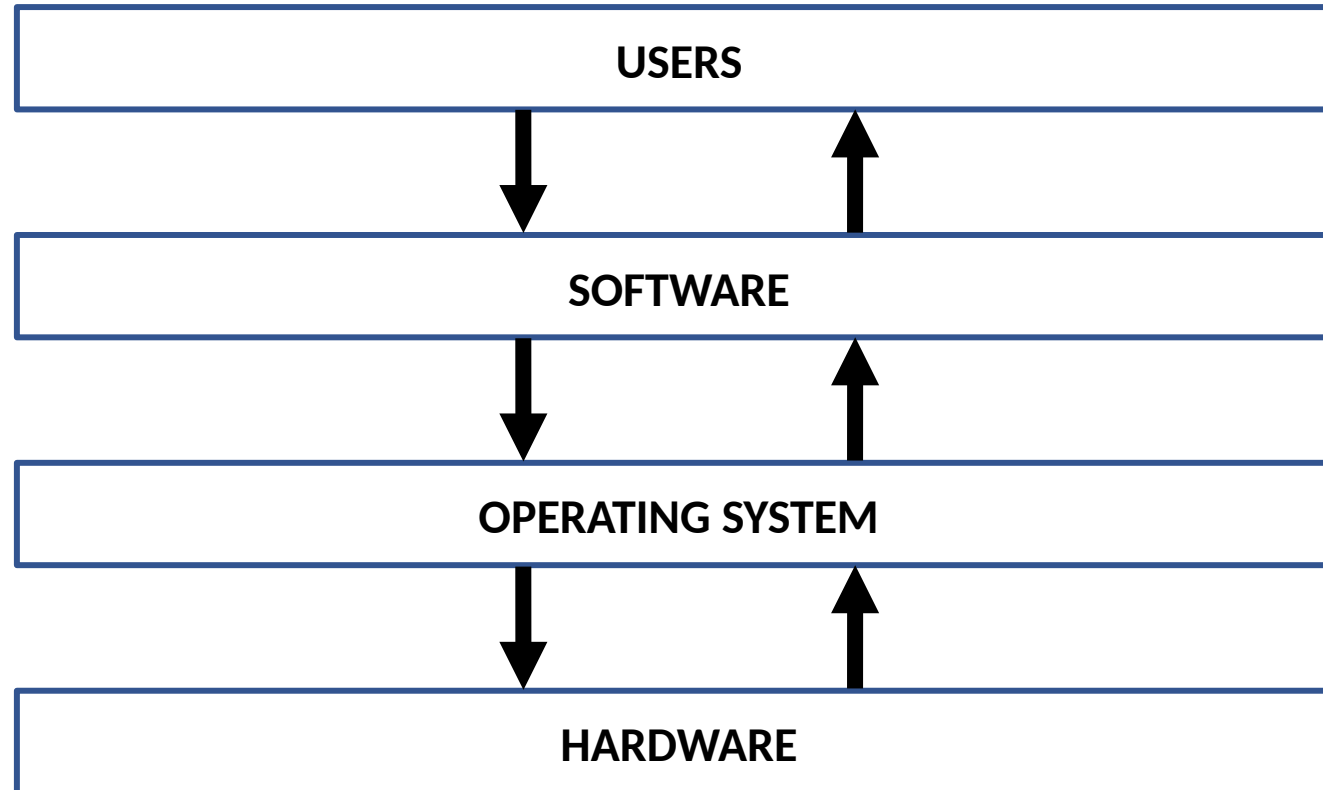
<https://opale.iut-clermont.uca.fr/info/wiki/doku.php?id=progsys:progsys>

Class Syllabus

1. Introduction
2. Process
3. File Management System
4. Pipelines
5. Signals
6. Memory Management
7. Ressources Management
8. Socket API using C Programming

Operating System (OS) : Definition

- Main program connecting the hardware and the software

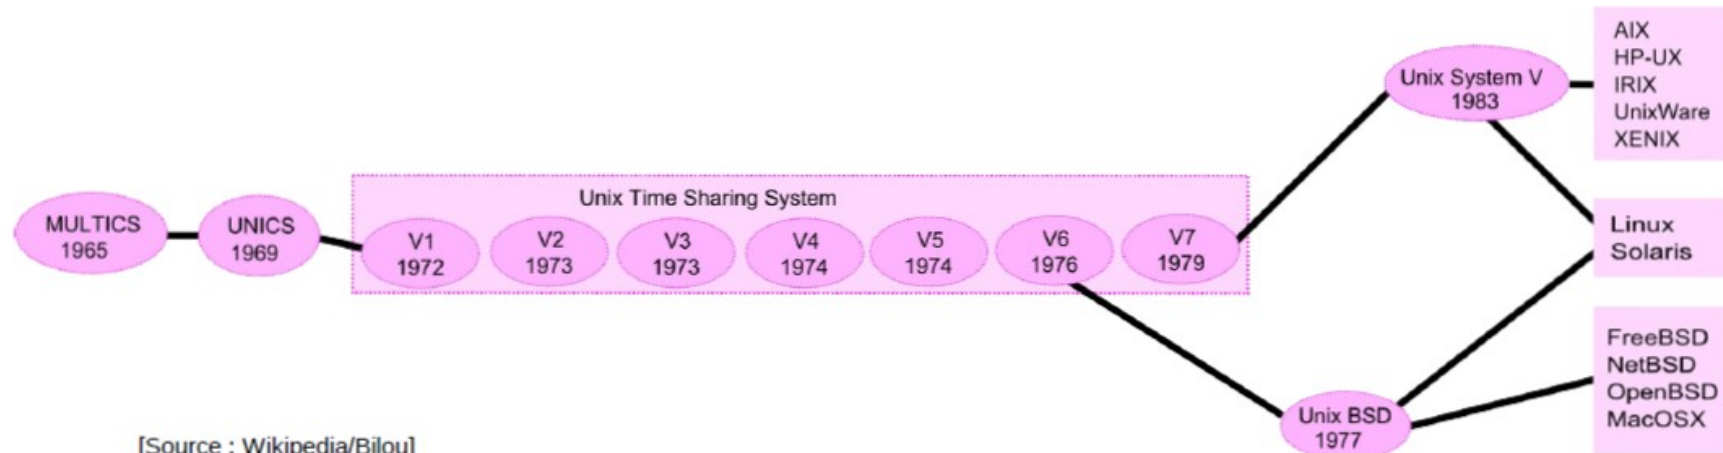


Operating System (OS) : Objectives

- Provide users with an interface masking the hardware
- Managing and sharing the resources of the machine
 - Process (order, communication)
 - Memory
 - File System
 - Network

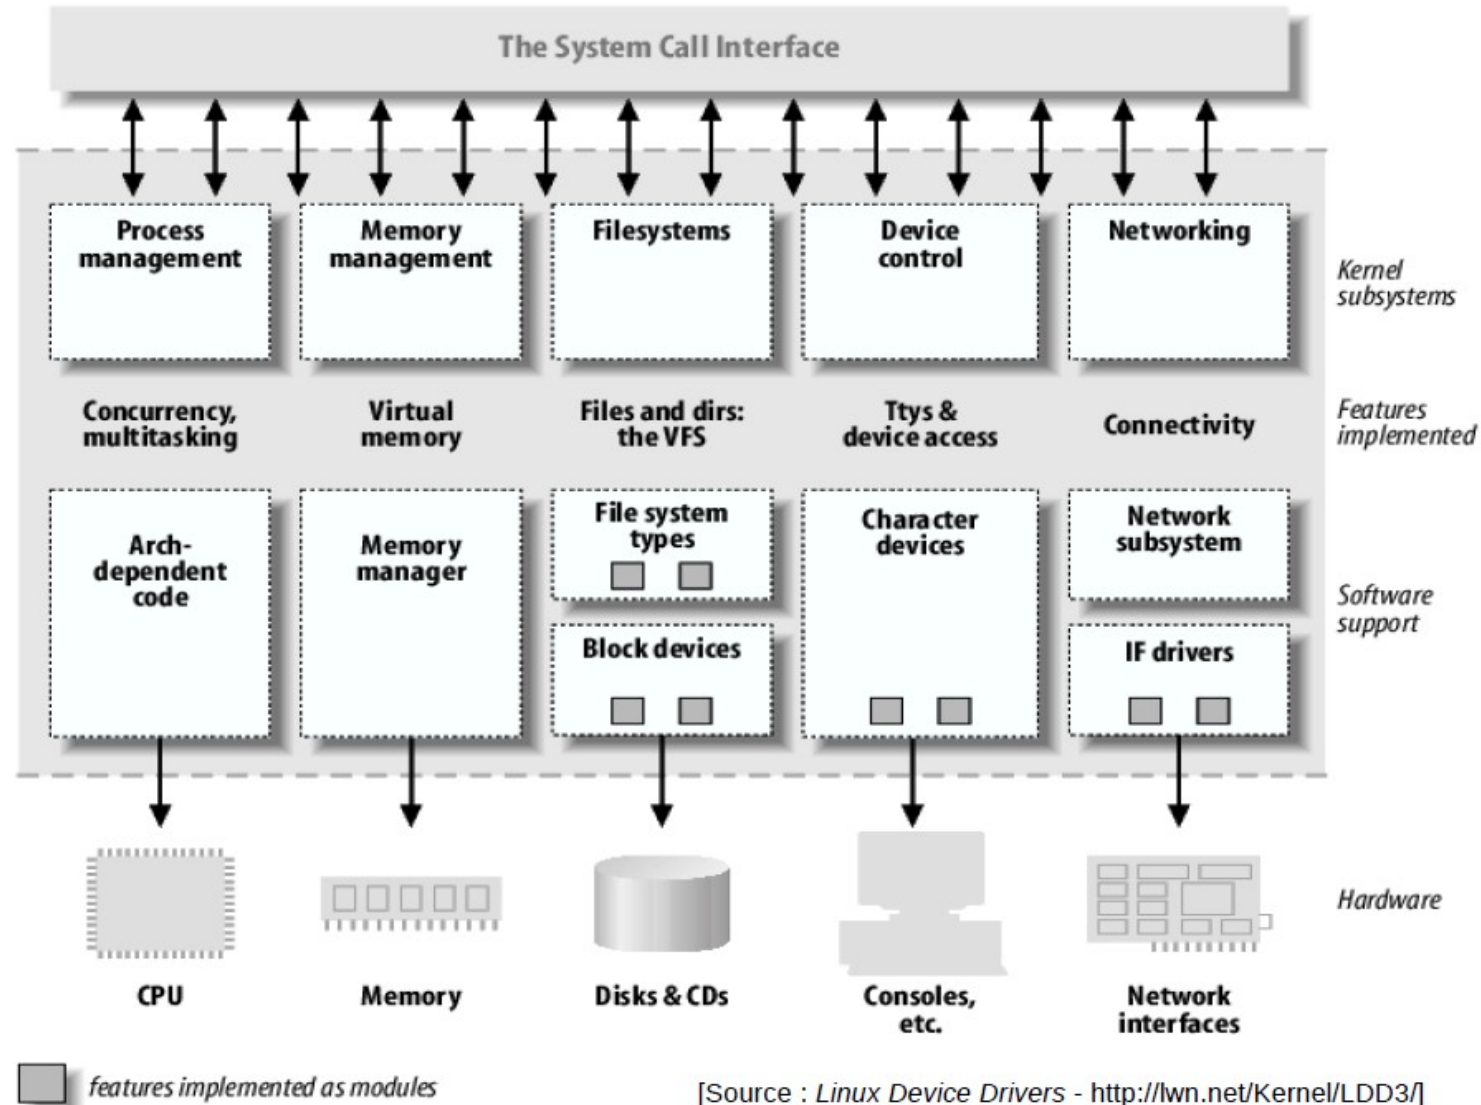
Operating System (OS) : Families

- DOS
- Windows
- UNIX →
 - AIX (IBM)
 - Solaris, Open Solaris
 - LynxOS (RTOS)
 - QNX (RTOS)
 - Linux
 - OpenBSD, FreeBSD
 - NetBSD
 - MacOS X
- AS400
- *etc.*



[Source : Wikipedia/Bilou]

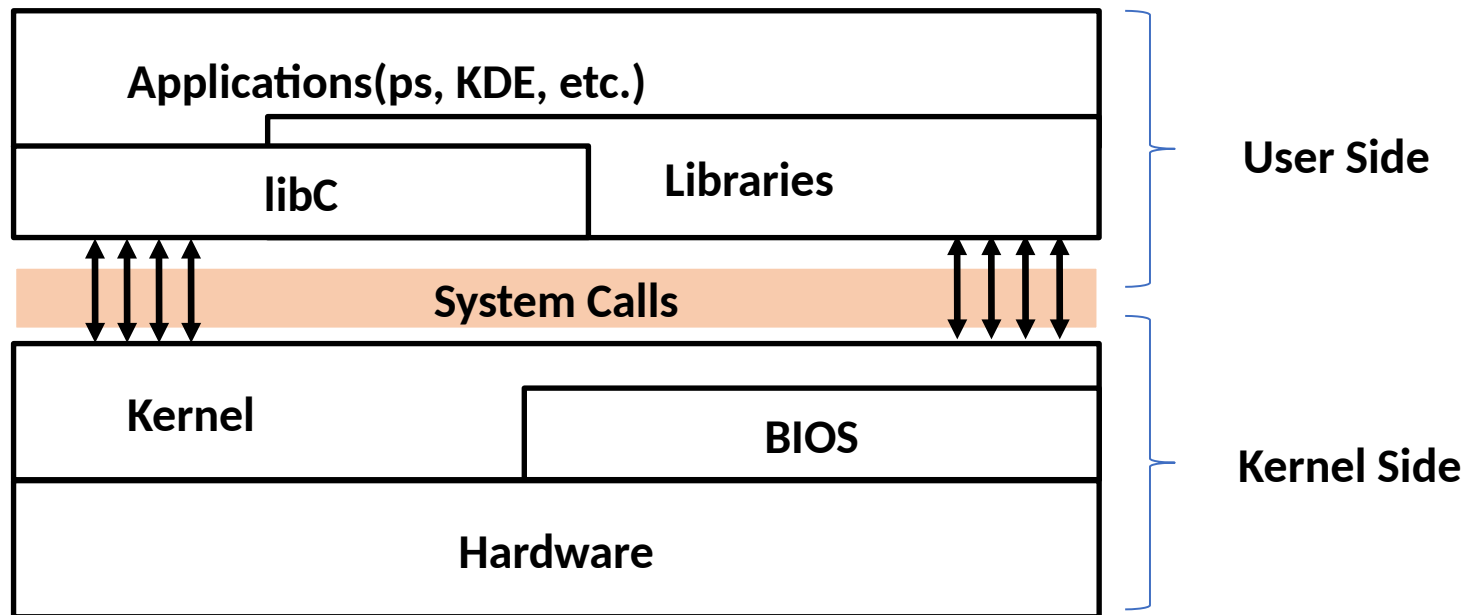
Linux Kernel Functions



[Source : *Linux Device Drivers* - <http://lwn.net/Kernel/LDD3/>]

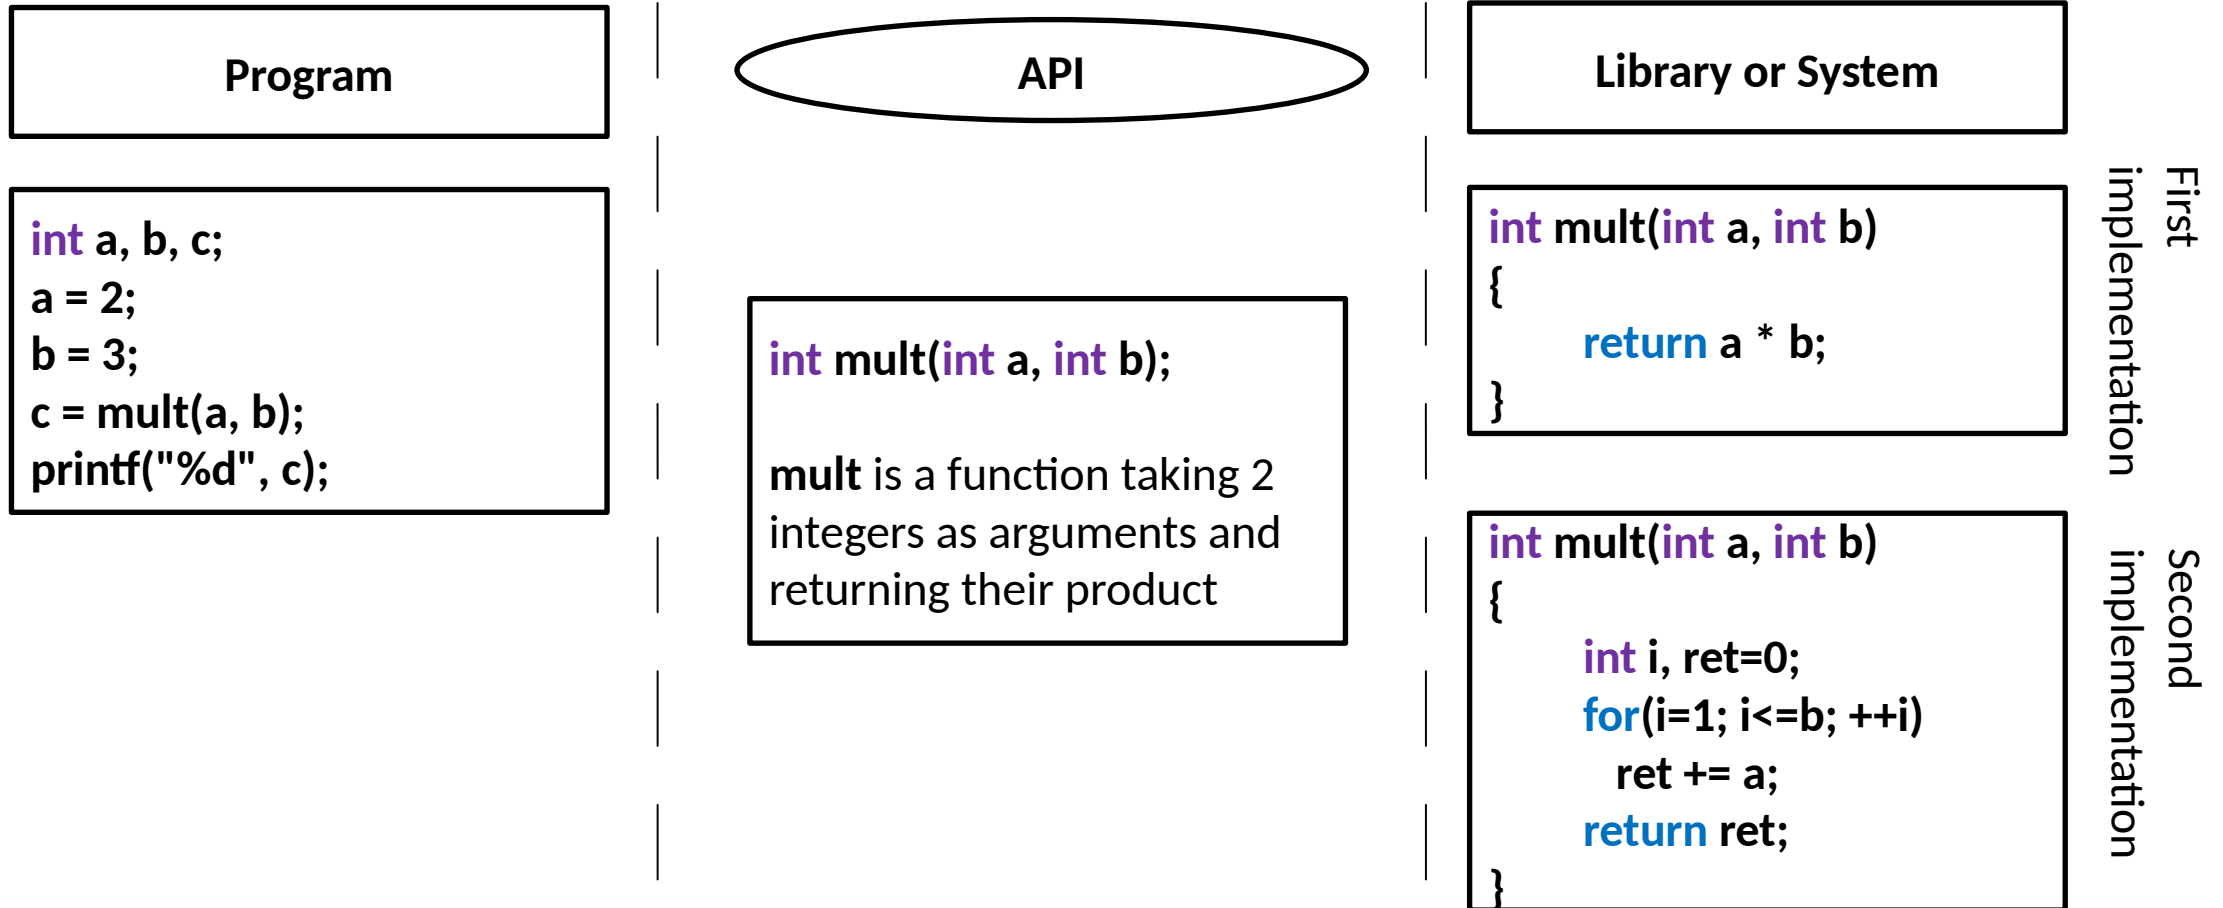
System Call (Syscall)

- Communication from the **Application layer** to the **System layer** is handled by the **System Call API**



Did I just read API?

- API = Application Programming Interface



"POSIX" Systems

- **POSIX = Portable Operating System Interface (Unix based)**
- Family of standards specified by the **IEEE (Institute of Electrical and Electronics Engineers)**
- **POSIX defines**
 - Shell Command-Lines (ksh, ls, man, etc.)
 - System call API
 - Real-time extensions
 - Thread API (light process)

Class Purpose

- The System Programming Class aims to present the various POSIX System Calls and how you can use them
- Those System Calls are regrouped in 6 main categories
 1. Process
 2. Hard Drive storage
 3. Pipes
 4. Signals
 5. Memory management
 6. Resources management
 7. Socket API

Chapter 2

Process

Class Syllabus

1.1) General Instructions

1.2) Cycle of life

1.2.1) Creation

1.2.2) Errors

1.2.3) Waiting

1.2.4) Replacement

1.2.5) Termination

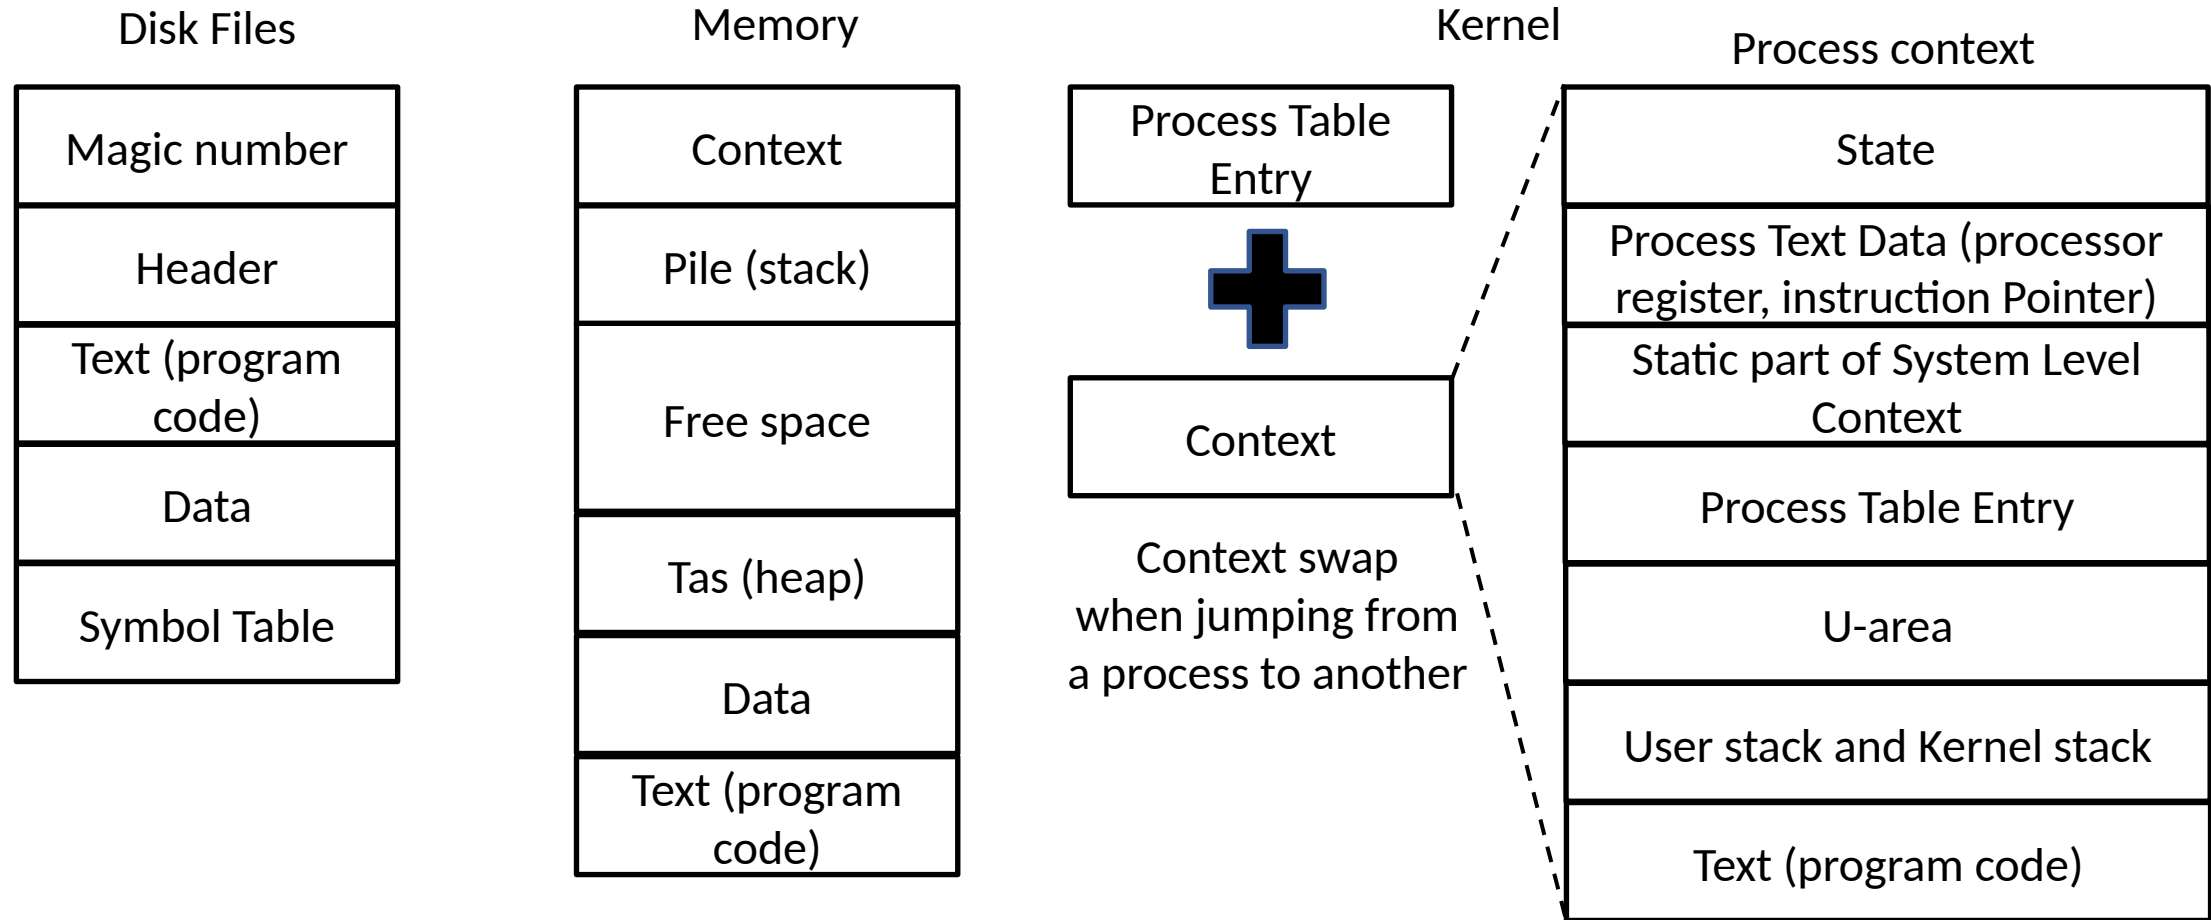
1.3) Communication between processes

1.4) Light Processes (threads)

Process : definition (1/2)

- A **process** is the running instance of a **program** (also know as **task**)
- A **program** is a **file** containing either programming code or data
- There is way more data needed to describe a process running into the system
- All this data allow the system to run multiple process at the same time and to jump from one to another

Process : definition (2/2)



Some of those data (Process Table, U-area) will be explained later

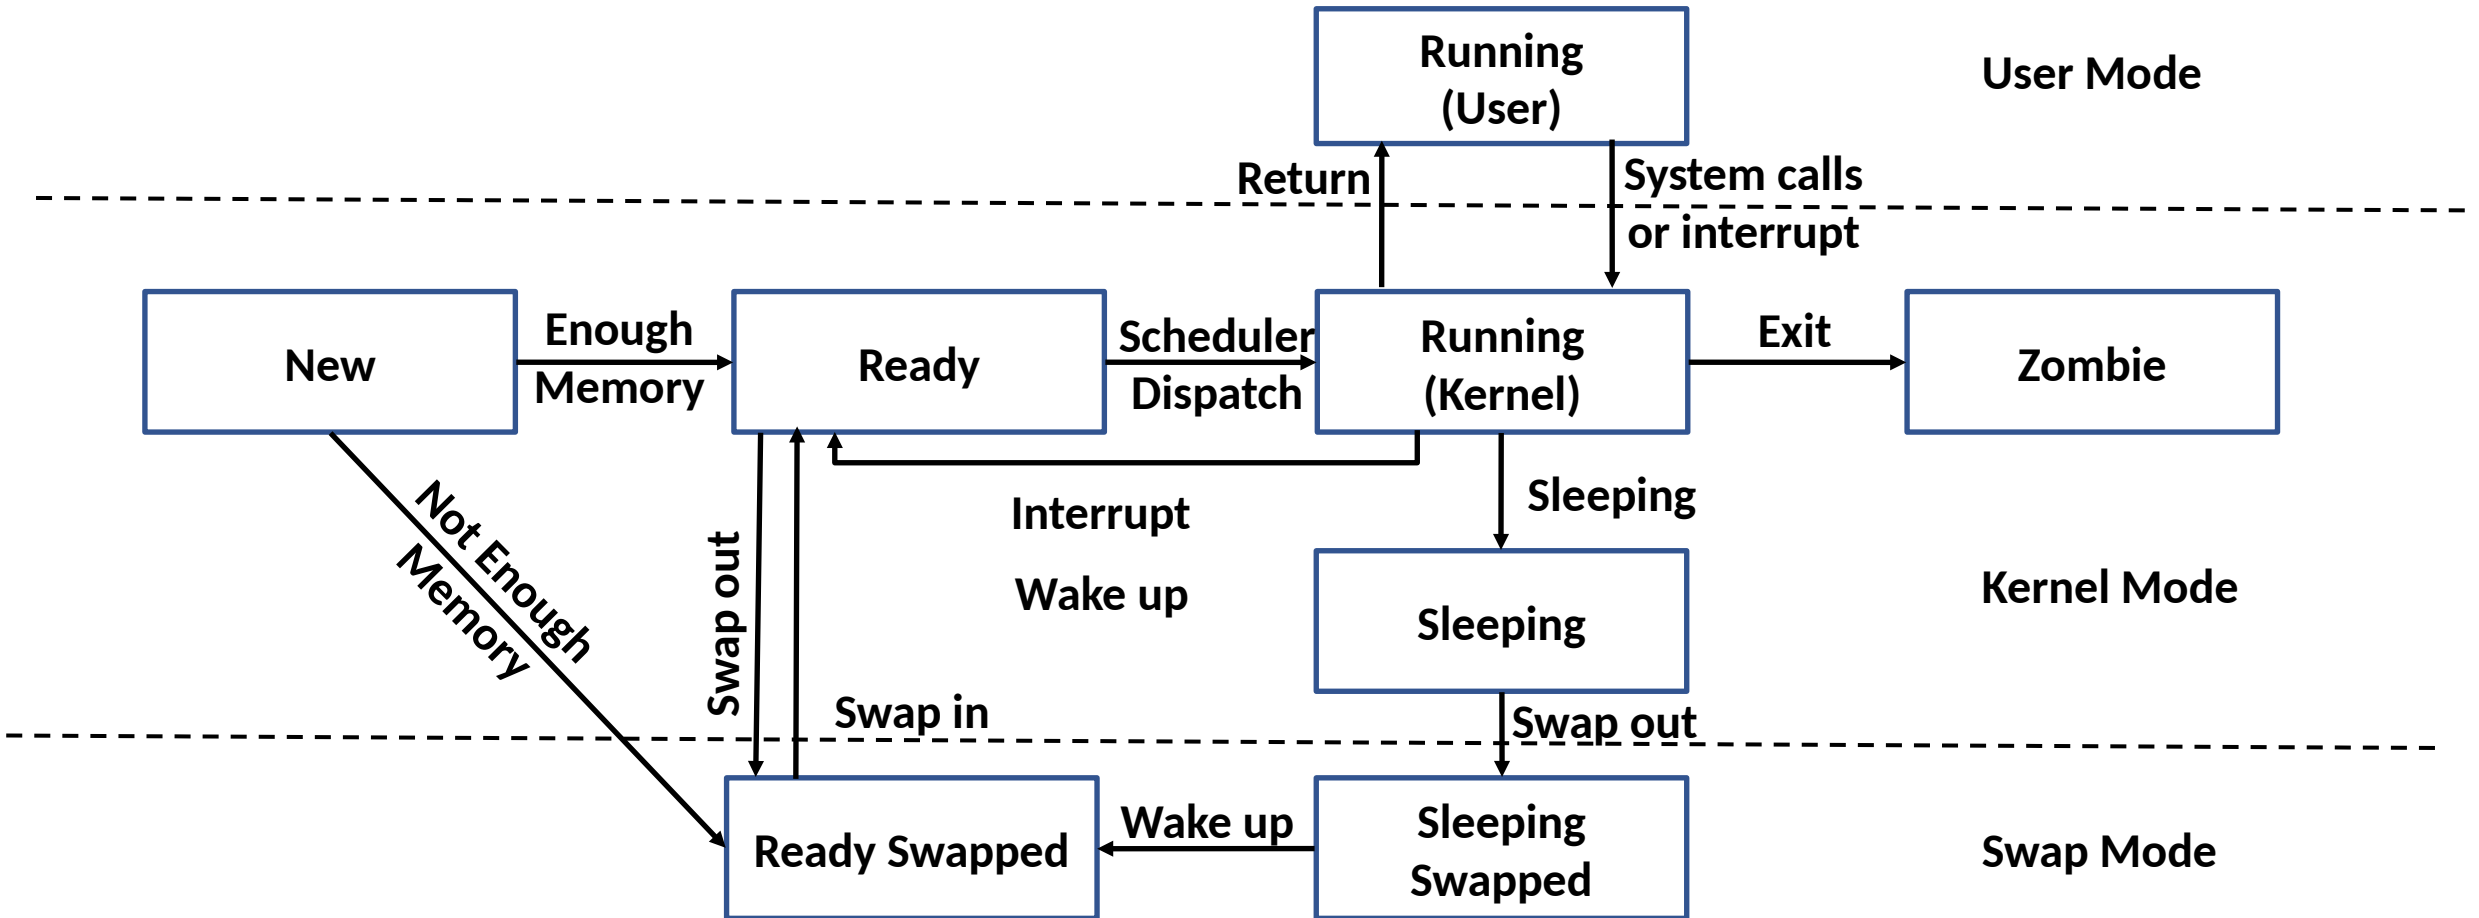
U-Area (User Area)

- Specific data of a process stored in a stack segment
 - Real and effective **UID (User ID)**
 - Timer field (user and system)
 - Reactions to signals (future classes)
 - Error field recording errors encountered during a system call
 - Return value field containing the result of system calls
 - Current directory (**pwd**) and current root (**chroot**)
 - User file descriptor table
 - Limits (man *ulimit*)
 - *umask* : masks mode settings on files the process creates

Process execution (1/2)

- The processor can only run one process at any time (except for a multiple cores processor)
- The System switch between processes, giving the illusion of multitasking
- Processes might be in different states depending on their current activity : new, ready, running, sleeping, terminated, etc.

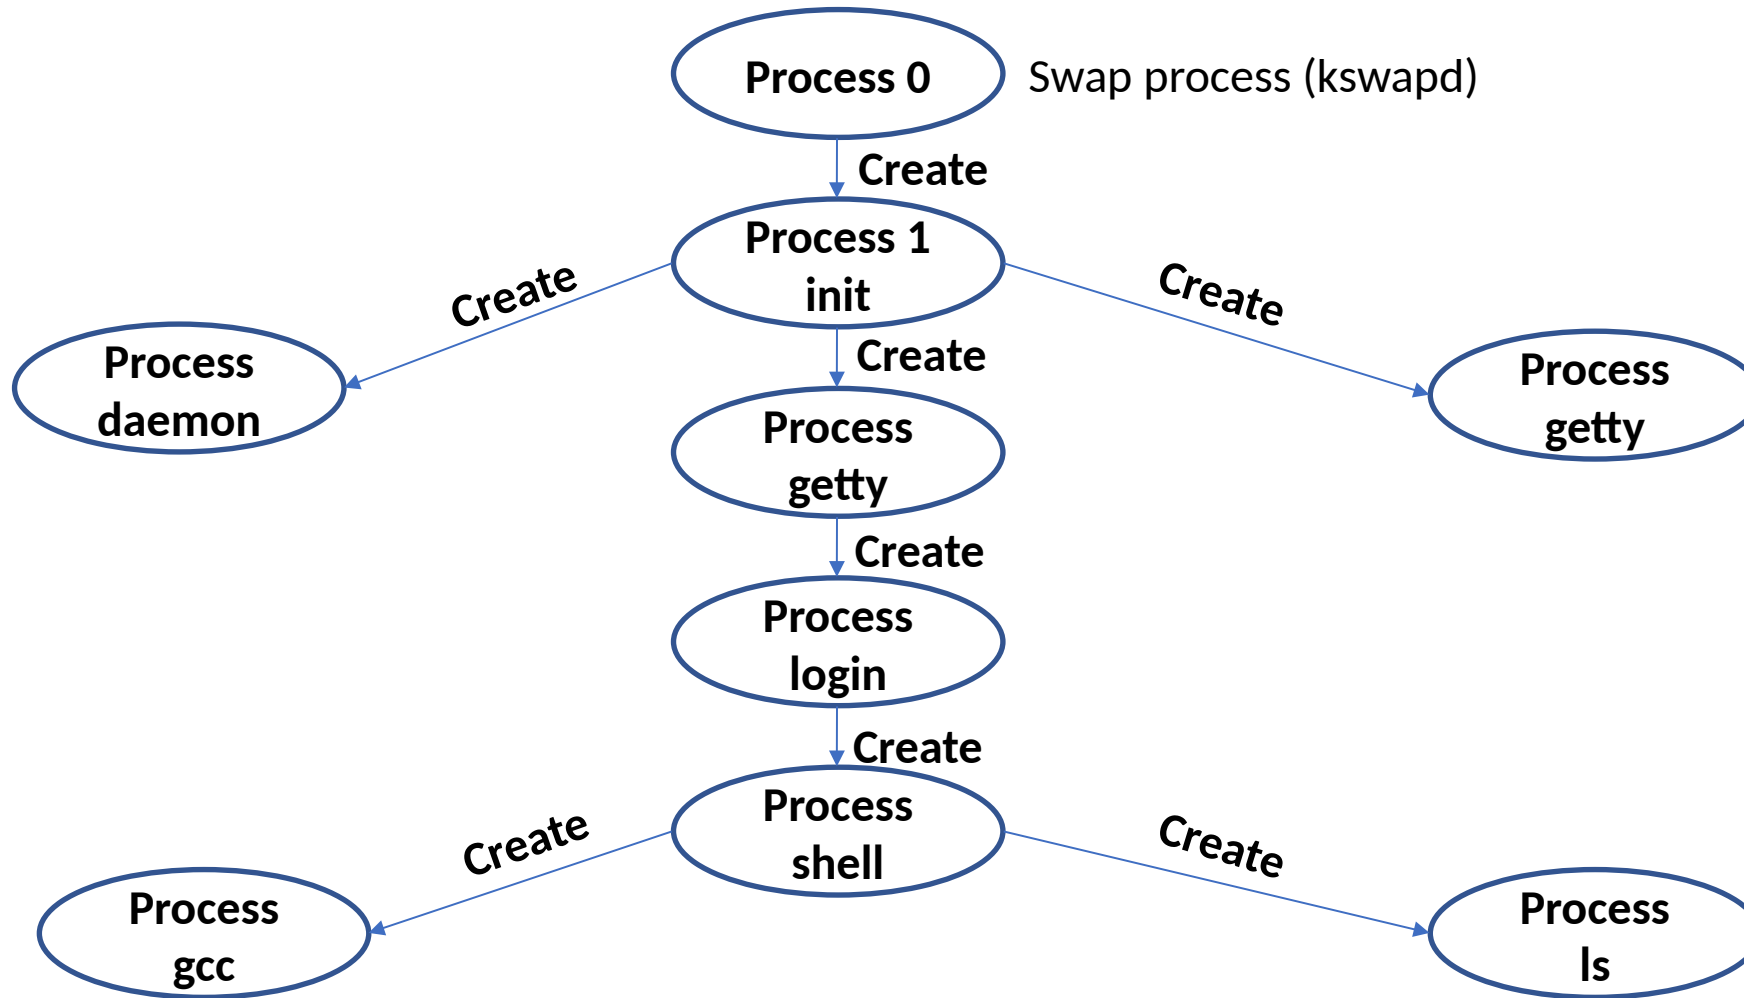
Process execution (2/2)



Process creation (1/4)

- All the processes share a **father/child** relationship : **process tree**
- Initial process is **0** (<swapper>). He gives birth to **1** (<init>)
- Each new process is the child of another, and any process can give birth to a child and become a father
- A father process **should** wait the death of all of his child before ending
- If a process die before his child, they will get adopted by <init> => no orphans

Process creation (2/4)



Process creation (3/4)

- Each process has his own **PID** (**P**rocess **ID**) and **PPID** (**P**arent **P**rocess **ID**)
- **PPID** is also a **PID**, both of **pid_t** type (integer)
- Using the "**ps**" command show the process tree

```
$> ps axo stat,ppid,pid,tty,user,comm
```

STAT	PPID	PID	TT	USER	COMMAND
S	0	1	?	root	init
S	1	2	?	root	keventd
Ss	1	816	?	root	inetd
Ss+	1	1119	tty1	root	getty
S	23094	6851	?	www-data	apache2
S+	14970	14974	pts/2	toto	gnuplot
R+	14956	17569	pts/1	toto	ps

Process creation (4/4)

- To create a new process, the **father** will duplicate himself and the new instance will load a new running code
- System primitives
 - **fork()** : create a new process by duplicating the caller
 - **exec()** : replace the programming code (in memory) of the clone by the one of the process to execute (read on hard disk) => replacement primitive
 - **wait()** : notify the father of the death of one of his child and enable the recovering of termination data (blocking call)
- **SIGCHLD** signal can also be sent by a child to inform his father

How to : `fork()` (1/3)

- `fork()` creates to duplicate of a process
- All the data of the process (U-area) are duplicated, with the exception of PID and PPID
- Opened file descriptors of the father are duplicated too => the child has the same open files
- Information about the child (running time, etc.) are reset

How to : fork() (2/3)

- Headers :

- `#include <sys/types.h>`

- `#include <unistd.h>`

- `pid_t` is of integer type

- `pid_t fork(void);`

- Other useful primitives :

- `pid_t getpid(void);`

- `/* My PID */`

- `pid_t getppid(void);`

- `/* PID of my father*/`

- `pid_t getcwd(char* buf, size_t size);`

- `/* What's my ...`

- `pid_t getwd(char* buf);`

- `... working directory */`

```
int main(int argc, char* argv[]) {  
    pid_t child = fork();    /* From here, father and child run the same programming code */  
    switch(child) { /* Differences starts here */  
        case -1 :  
            perror("fork");  
            exit(errno);  
        case 0:    /* Child programming code */  
            printf("Child : my PID is %d\n", getpid());  
            printf("Child : my father's PID is %d\n", getppid());  
            break;  
        default: /* Father programming code */  
            printf("Father : my child's PID is %d\n", child);  
    }  
    printf("%d : This sentence will appear in both processes\n", getpid());  
    return 0;  
}
```


Handling system errors (1/2)

- Each system function returns a return code
 - **fork()** return -1 if an error occurs
 - **getenv()** return NULL if an error occurs
- Return code doesn't tell the type of error encountered
- How to know the type of this error
 - Global variable **errno**
 - Defined in **<errno.h>**
 - **errno** value isn't significant if there's no error encountered
 - **errno** value is specific for each system call (mentioned in *man*)

Handling system errors (2/2)

- Display the error message
 - **void perror (const char *msg)**
- **perror()** display a message on the standard error output with the last know error encountered by a system call (see *man*)
- In System Programming (like every other type of programming), error handling is **MANDATORY!**

Using strace

- **Strace** shell command enable you to follow every system call (and signal)

Your very best friend in SysProg ◀◀

- **Usage exemple :**

```
$> strace cat /dev/null
[...]  
open("/dev/null", O_RDONLY|O_LARGEFILE) = 3 [...]  
$> strace cat /dev/lapinblanc  
[...] open("/dev/lapinblanc", O_RDONLY|O_LARGEFILE)=-1 ENOENT (No such file or  
directory)  
[...]
```

How to : wait (1/3)

- After a **fork()**, the father can use the **wait()** function
- **wait()** is a blocking function, waiting for any child of the father
- **waitpid()** is a blocking function, waiting for any or a specified child

```
#include <sys/types.h>
#include <sys/types.h>

/* Waiting for a child death, informations recovered in status */
pid_t wait(int* status);

/* Waiting for a specified child death (or any) */
pid_t waitpid(pid_t wpid, int* status, int options);
```

How to : wait (2/3)

- Interpretation of **status** is done using macros
 - **WEXITSTATUS**, **WCOREDUMP**, etc. (man **2** wait)

```
/* [...] in the father, after using a fork() */
```

```
int status = 0;
```

```
pid_t pid;
```

```
if ( ( pid = wait(&status) ) == -1 ) {  
    perror("wait");  
    exit(errno);  
}
```

```
printf("My son %d ended with the exit code %d\n", pid, WEXITSTATUS(status));
```

```
/* [...] */
```

```
int main(int argc, char* argv[]) {  
    int status = 0;  
    pid_t returnCode;  
    pid_t child = fork();  
    switch(child) {  
        case -1 :  
            perror("fork"): exit(errno);  
        case 0:    /* Child programming code */  
            printf("Child : my PID is %d\n", getpid());  
            break;  
        default:  /* Father programming code */  
            printf("Father : my child's PID is %d\n", child);  
            returnCode = wait(&status);  
            if ( returnCode = -1 )  
                printf("My child %d ended with the code %d\n", returnCode, WEXITSTATUS(status));  
    }  
    return 0;  
}
```

How to : exec (1/2)

- **exec()** is replacement primitive, replacing the running code of a process by another one
- **exec()** launch a process with arguments, like while using the command line
- **exec()** family enables to specify the arguments of the child in different ways

```
#include <unistd.h>
/* exec() family */
int execl(const char *path, const char *arg, ..., NULL);
int execlp(const char *file, const char *arg, ..., NULL);
int execl(const char *path, const char *arg, ..., NULL, char *const *envp[]);
int execv(const char *path, const char *argv[]);
int execvp(const char *file, const char *argv[]);
int execvp(const char *file, const char *search_path, const char *argv[]);
```

2 – Process > 1.2 Cycle of life > 1.2.4 Replacement

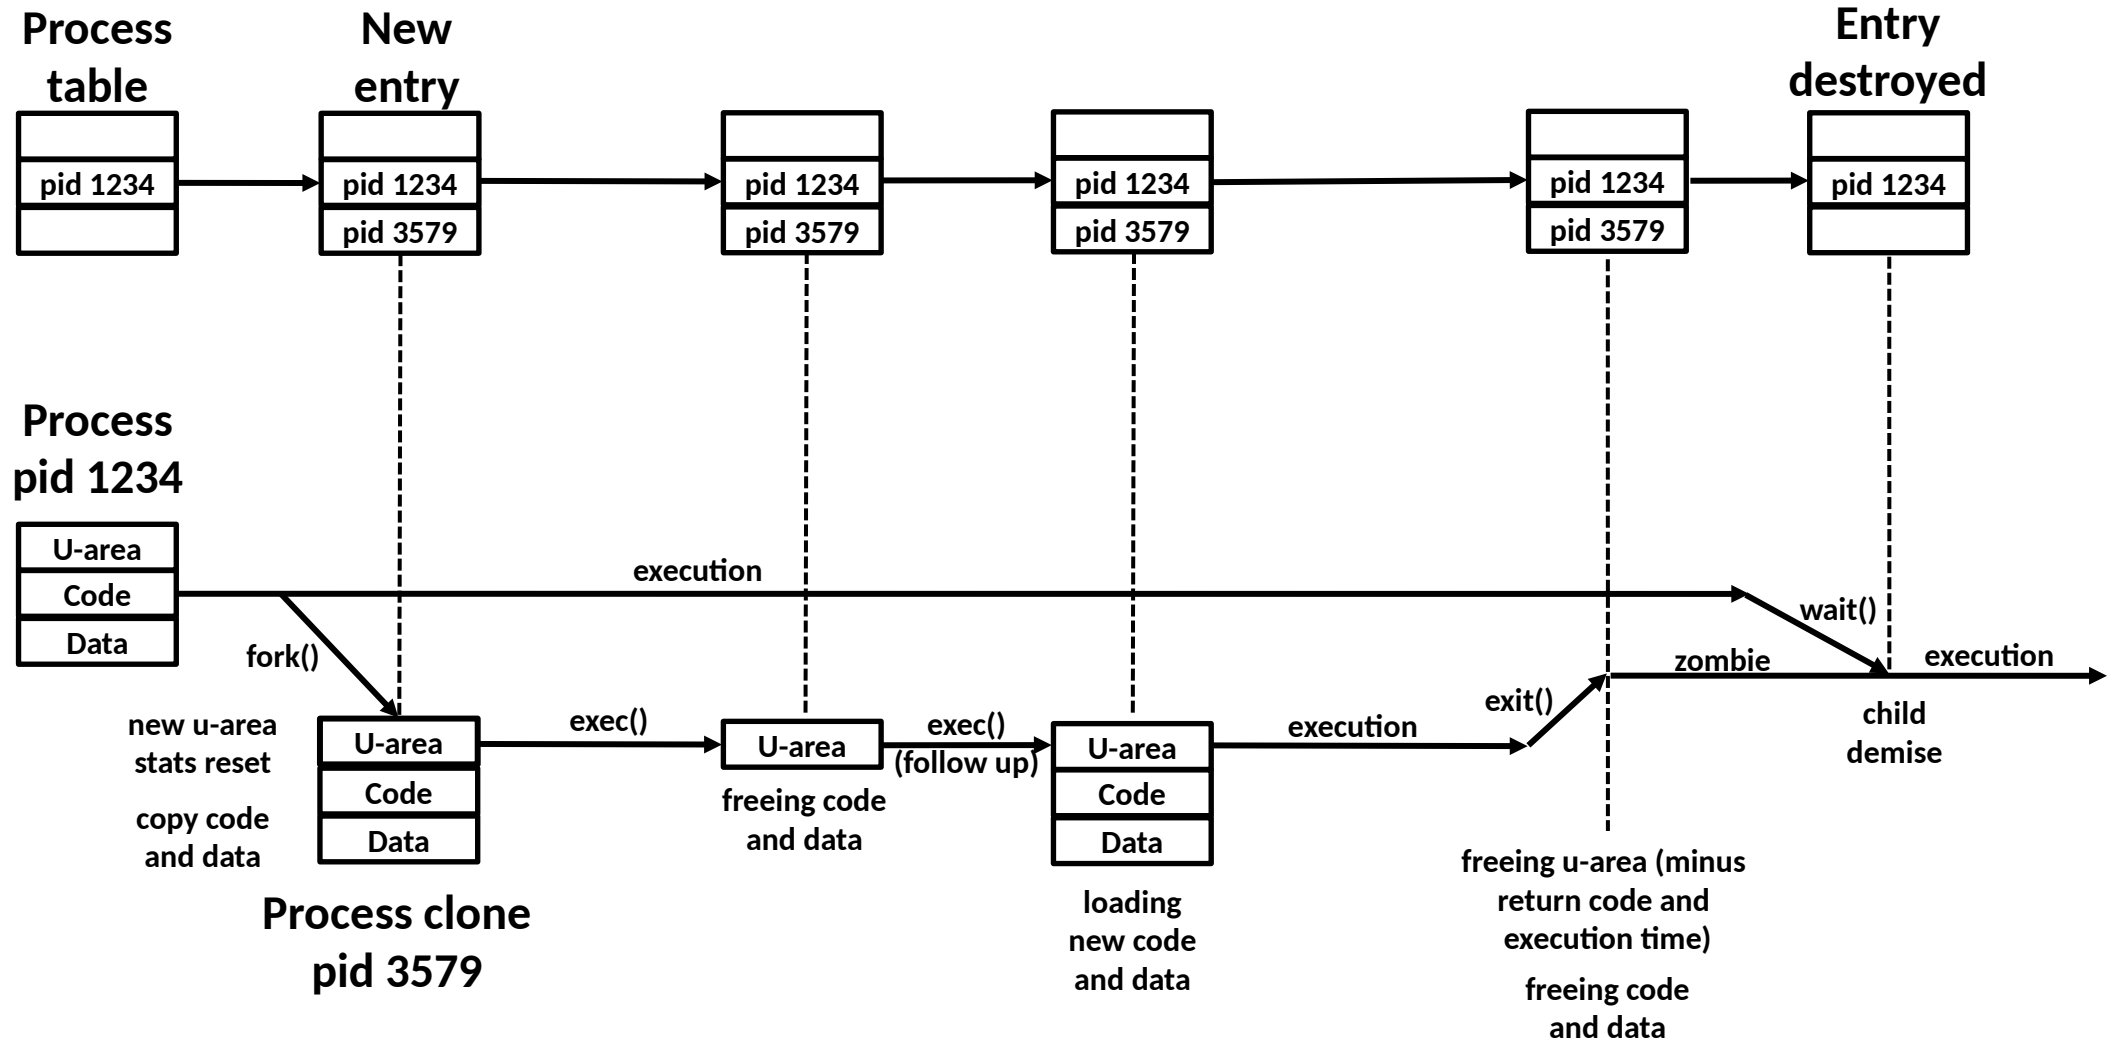
How to : exec (2/2)

```
int main(int argc, char* argv[]) {
    int status = 0, error = 0;
    pid_t returnCode;
    pid_t child = fork();
    switch(child) {
        case -1 :
            perror("fork"): exit(errno);
        case 0:    /* Child programming code */
            printf("Replacement of the code, here we will do a <ls -l>\n");
            error = execl("/bin/ls", "ls", "-l", NULL);
            printf("This will never appear, except if the execlp call has failed\n");
            fprintf(stderr, "execlp() failed with the error code %d\n", error);
            exit(errno); break;
        default:  /* Father programming code */
            printf("Father : my child's PID is %d\n", child);
            if ( wait(&status) = -1 ) { perror("wait"); exit(errno); }
    }
    return 0;
}
```


Processes Death

- A process doesn't immediately disappear at the end of his running time
- In fact, when a process is terminated (call to **exit()**), he becomes a zombie
- He will disappear only after the father hears about his termination (using **wait()** or a signal)
- *<init>* is always in **wait()**, which enables him to adopt orphans processes so they can properly disappear

Process life: Summary



Sharing data between father and child

- **Data** (variables) are duplicated (exception : open file descriptors)
- **Memory** isn't shared between father and child => modifications in a process aren't visible or reverberated in the other one
- Duplication is done using *copy-on-write* (optimisation non négligeable)
- To use shared memory, see the system API (**shm_open()**, **shm_unlink()**, etc.)

Exec() arguments and line command : argc and argv (1/2)

- `int` main(`int` argc, `char`* argv[]) : why?
- `argc` : number of line command arguments
- `argv` : the arguments, in an array of `char`
- `argv[0]` : name of the program (`argc` is always ≥ 1)

```
int main(int argc, char* argv[])
{
    int i = 0;
    for (i = 0; i < argc; ++i)
        printf("argument %d : %s\n", i, argv[i]);
    return 0;
}
```

```
/home/prof> ./toto -a1F ab 12
argument 0 : ./toto
argument 1 : -a1F
argument 2 : ab
argument 3 : 12
/home/prof>
```

Exec() arguments and line command : argc and argv (2/2)

- **exec()** family uses arguments : number of line command arguments

```
pid_t child = fork();
switch(child) {
    case -1 :
        perror("fork"); exit(errno);
    case 0:    /* Child programming code */
        printf("Replacement of the code, here we will do a <ls -l>\n");
        error = execl("/bin/ls", "ls", "-l", NULL);
        printf("This will never appear, except if the execlp call has failed\n");
        fprintf(stderr, "execlp() failed with the error code %d\n", error);
        exit(errno);    break;
    default:  /* Father programming code */
        printf("Father : my child's PID is %d\n", child);
        if ( wait(&status) = -1 ) { perror("wait"); exit(errno); }
}
```

Using the environment (1/2)

- The environment is built on key/values associations
 - *Environment variable* name, **char** value
- Examples : **PS1, USER, SHELL, etc.**
- Father's environment is inherited by any child

```
#include <stdlib.h>
char* getenv(const char *name);           /* return the value of a key */

char* setenv(const char *name, const char *value, int overwrite); /* define a key/value pair */

/* take a string "key=value" and execute setenv("key", "value", "1") */
int putenv(const char *string);

void unsetenv(const char *name);          /* delete a key (and also the value) */
```

Using the environment (2/2)

```
int main(int argc, char* argv[]) {
    int status = 0;
    pid_t returnCode;
    setenv("toto", "titi", 1);          /* changing the environment before the fork */
    pid_t child = fork();
    switch(child) {
        case -1 :
            perror("fork"): exit(errno);
        case 0:    /* Child programming code */
            printf("toto = %s\n", getenv("toto")); /* display "titi" */
            break;
        default:  /* Father programming code */
            if ( wait(&status) = -1 ) { perror("wait"); exit(errno); }
    }
    return 0;
}
```

Sharing data between father and child

- **fork()** creates a new concurrent process
- A process can also split himself in several **threads**, contained inside the same process
- **Threads** runs in **concurrent computing** but share their resources => allow the developer to make a implement a bit of parallelism (**see you in 3A**)
- **Threads** use API like **POSIX Threads**

POSIX Threads

- **<pthread.h>**
- **pthread_create** create a new thread to run a specific function
- A thread ends with the execution of **pthread_exit()** OR when it reaches the end of the running function
- Synchronisation (waiting) : **pthread_join()**

2 – Process > 1.4 Light Processes (threads)

```
void* thread_function(void* data) {
    int* x = (int*) data;
    printf("Value %d\n", *x);
    *x = 4;    /* changing value (no possible conflict here). thread stops here, same as pthread_exit() */
    return 0;
}

int main(int argc, char* argv []) {
    int value = 0, error = pthread_create(&t, NULL, &fonction_du_thread, &valeur);
    pthread_t t;
    if (!error) {
        printf("thread creates, waiting for the end\n"); error = pthread_join(t, NULL);
        if (error) {
            fprintf(stderr, "pthread_join : %s\n", strerror(error)); exit(error);
        }
        printf("value should be equal to 4 : %d\n", valeur); return 0;
    }
}
```