

Chapter 3

Pipelines

Class Syllabus

4.1) Basic knowledge

4.2) Anonymous Pipes

- pipe()
- Read
- Write
- Read/Write
- Example

4.3) Named Pipes

4.4) Duplication

- dup()
- dup2()
- Usage

4.5) Exercise

Basic knowledge (1/2)

- Pipelines (or pipes) are:
 - A communication mechanism
 - Using characters passing
 - Transferring data between 2 local processes
- Pipes are FIFO structures:
 - Characters order in input is the same for the output
 - The reading action remove and destroy the content

Basic knowledge (2/2)



- Default behavior: reading from an empty pipe is blocked
- Default behavior: writing into a full pipe is blocked
- Pipes have a definite size of some Koctets (4KiB in Linux systems)

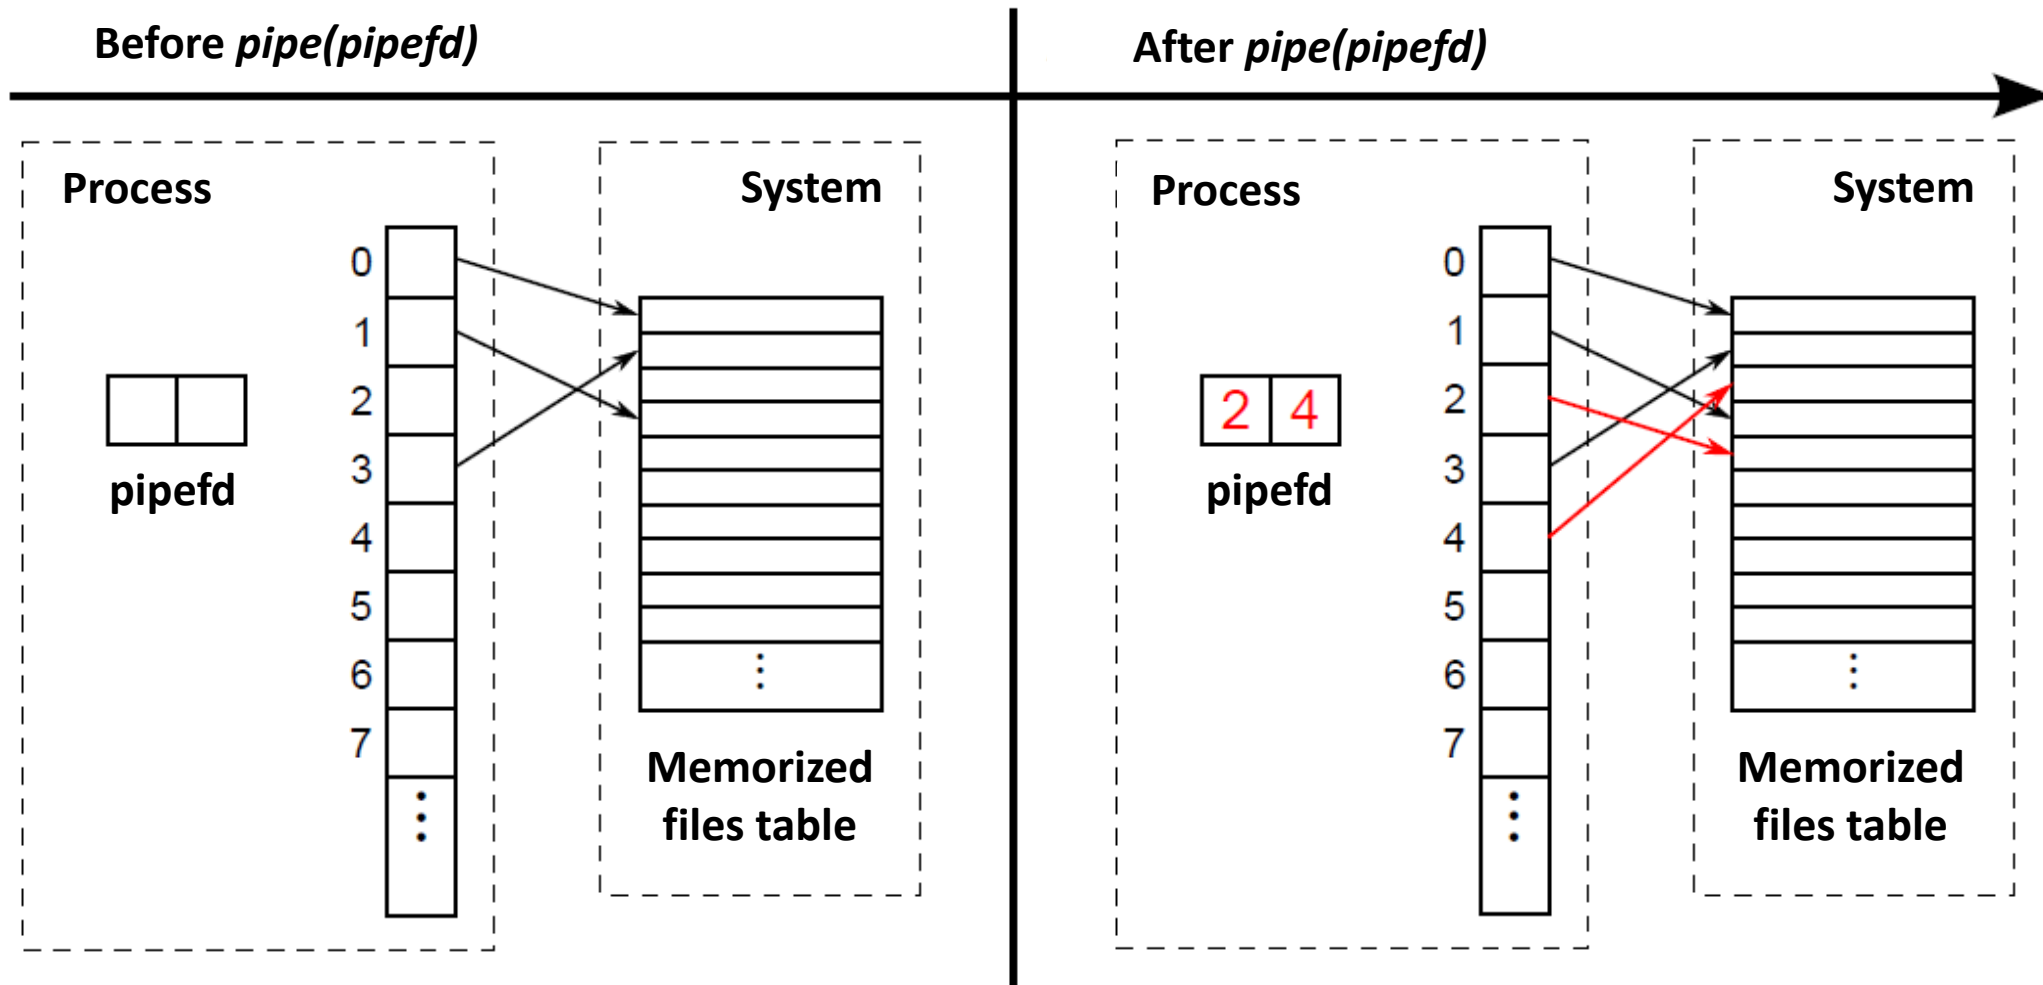
Anonymous Pipes

- A process can only use pipes:
 - he created by himself (with *pipe()*)
 - he herited from his father (descriptors heritage with *fork()* and *exec()*)
- Generally speaking, 2 processes (created using *fork()*) share the pipe
- They use the *read()* and *write()* system calls to send data to each other
- Also known as *tubes volatiles*

pipe() system call (1/2)

- `#include <unistd.h>`
- `#int pipe(int pipeFileDescriptor[2]);`
- *pipe()* create a pipe and a pair of file descriptors
- Each of those descriptors refer to an inode of a tube (one for each end)
- *pipe()* assign the 2 descriptors into an array:
 - `pipeFileDescriptor[0]`: reading descriptor
 - `pipeFileDescriptor[1]`: writing descriptor

pipe() system call (2/2)



Read(from a pipe) (1/2)

- pipeFileDescriptor[0] is the reserved reading descriptor
- Reading from a pipe is done using the system call *read()*

```
#define SIZE_BUFFER 1024
```

```
char buffer[SIZE_BUFFER];
```

```
char nbRead;
```

```
nbRead = read(pipeFileDescriptor[0], buffer, SIZE_BUFFER);
```


Read(from a pipe) (2/2)

read() system call behavior

if *pipe* is *not empty* **AND** contains size characters **then**

Reading nbRead = min(size, SIZE_BUFFER) characters

else if *pipe* is *empty* **then**

if *number of writers* is *null* **then**

end of the file => nbRead == 0

else if *number of writers* is *not null* **then**

if reading is *blocked* **then**

sleeping

else if reading is *not blocked* **then**

according to *indicator* **do**

case **O_NONBLOCK**: nbRead==-1 and errno==EAGAIN

case **O_NDELAY**: nbRead==0

Write(from a pipe) (1/2)

- pipeFileDescriptor[1] is the reserved writing descriptor
- Writing into a pipe is done using the system call *write()*

```
#define SIZE_BUFFER 1024
myType buffer[SIZE_BUFFER];
int nbWrite;
int n;          /* n < (SIZE_BUFFER * sizeof(myType)) */
nbRead = write(pipeFileDescriptor[1], buffer, n);
```
- Writing is an atomic operation if the number of characters to write is inferior to PIPE_BUFFER, size of a pipe for the system (see <limits.h>)

Write(from a pipe) (2/2)

`write()` system call behavior

if *number of readers* is *null* **AND** then

Sending *SIGPIPE* signal to the writer

else if *number of readers* is *not null* then

if *writing* is *blocked* then

`write()` will make a return only when the *n* characters are written into the pipe

else if *writing* is *not blocked* then

if *n* > PIPE_BUF then

return a number < *n*, maybe -1

else if *n* <= PIPE_BUF then

if *n free spaces* then

writing of *n* characters into the pipe (and *nbWrite=n*)

else

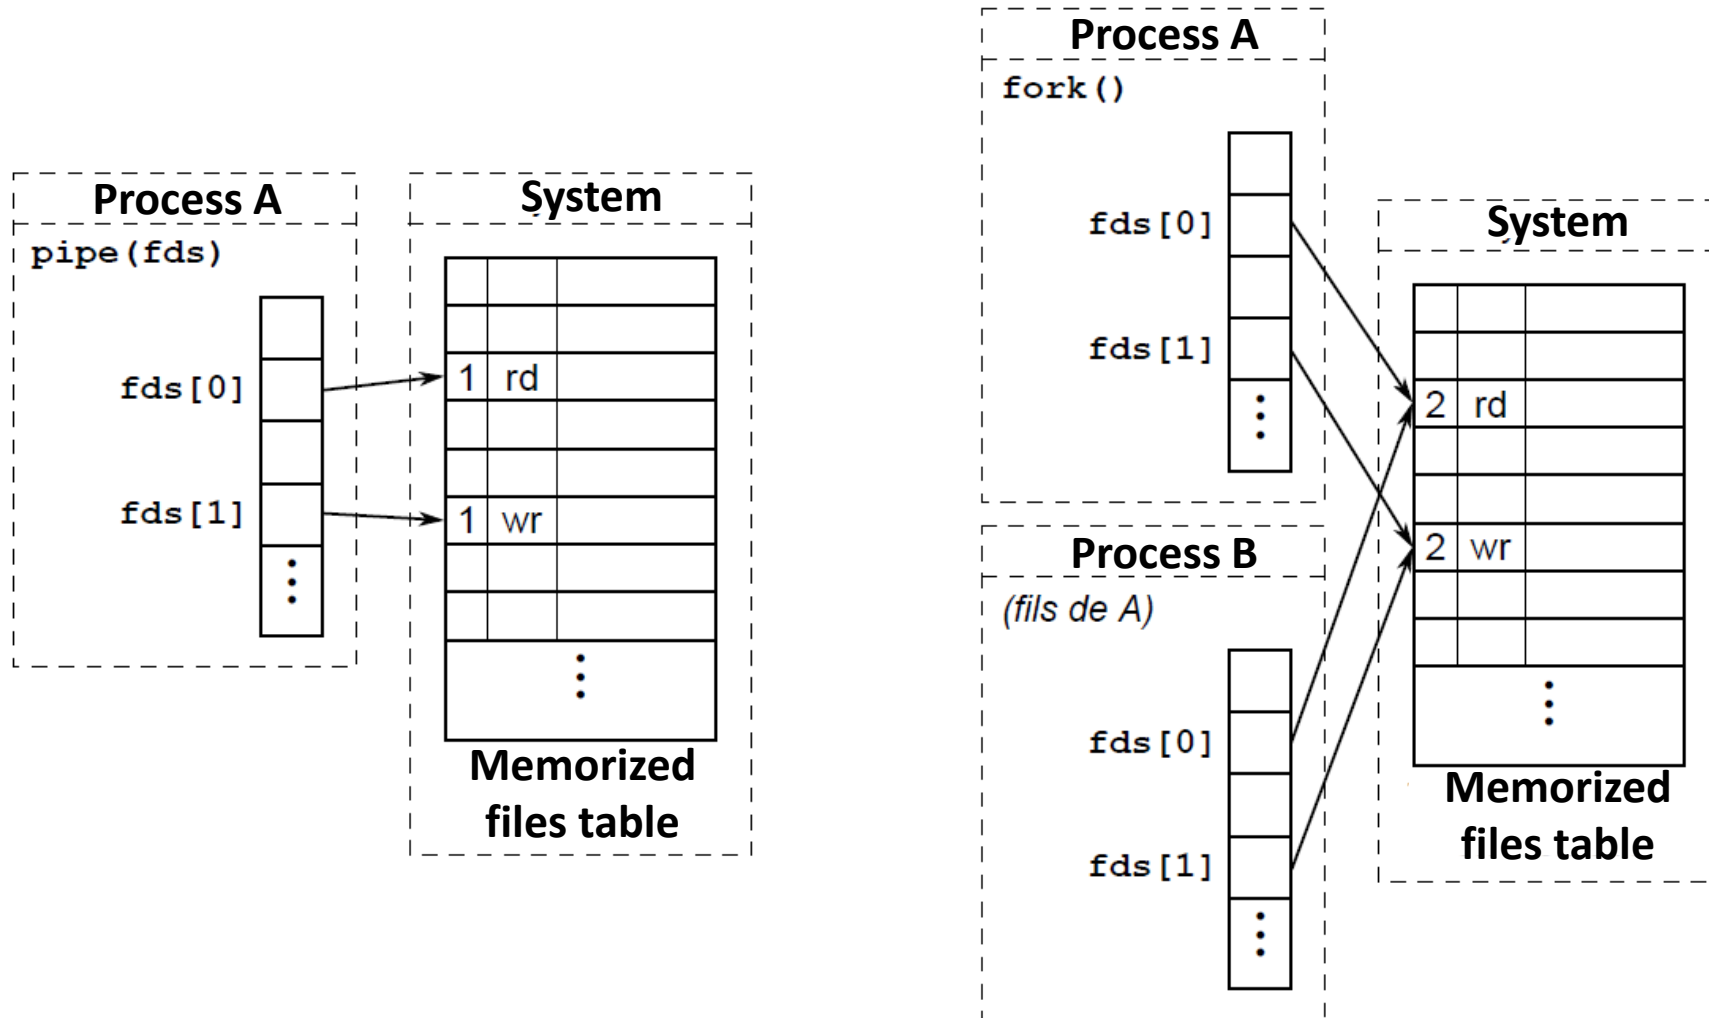
return -1 or 0

Read from/Write into a pipe (1/2)

- **Notable Information**

- The previous reading algorithm (managing writing) depends of the number of writers (managing reading). This number **MUST BE** up to date at anytime
- Any process owning a descriptor on an unused pipe entry **MUST** close it (with *close()*)
- Assuming the pipe descriptors can only be obtained using the system call *pipe()* or through heritage
 - There is an obligatory parental relationship between 2 processes communicating using a pipe

Read from/Write into a pipe (2/2)



Example of code with pipe() usage

```
void childCode(int pipe[2]) {  
    int d;  
    read(pipe[0], &d, sizeof(int));  
    printf("Child: reading %d\n", d);  
    close(tube[0]);  
    exit(0);  
}
```

```
void fatherCode(int pipe[2]) {  
    int e = 10;  
    write(pipe[1], &e, sizeof(int));  
    close(tube[1]);  
    wait(NULL);  
    exit(0);  
}
```

Named Pipes (1/3)

- Named Pipes are referenced pipes in the file system (name, etc.)
- Creation: `int mkfifo(const char *pathname, mode_t mode);`
- The resulting named piped will exist as long as:
 - his file system entry is not deleted**AND**
 - the number of processes that opened him is not null

Named Pipes (2/3)

Notable information:

- Considering a named pipe is open using his registered name in the file system, there is no parental relationship allowed between the processes using it
- Unlike an anonymous one, a named pipe must be opened with the *open()* system call before usage
- Opening a named pipe is done **EXCLUSIVELY** in *O_RDONLY* or *O_WRONLY* to assure the system can count the number of writers and readers at all time

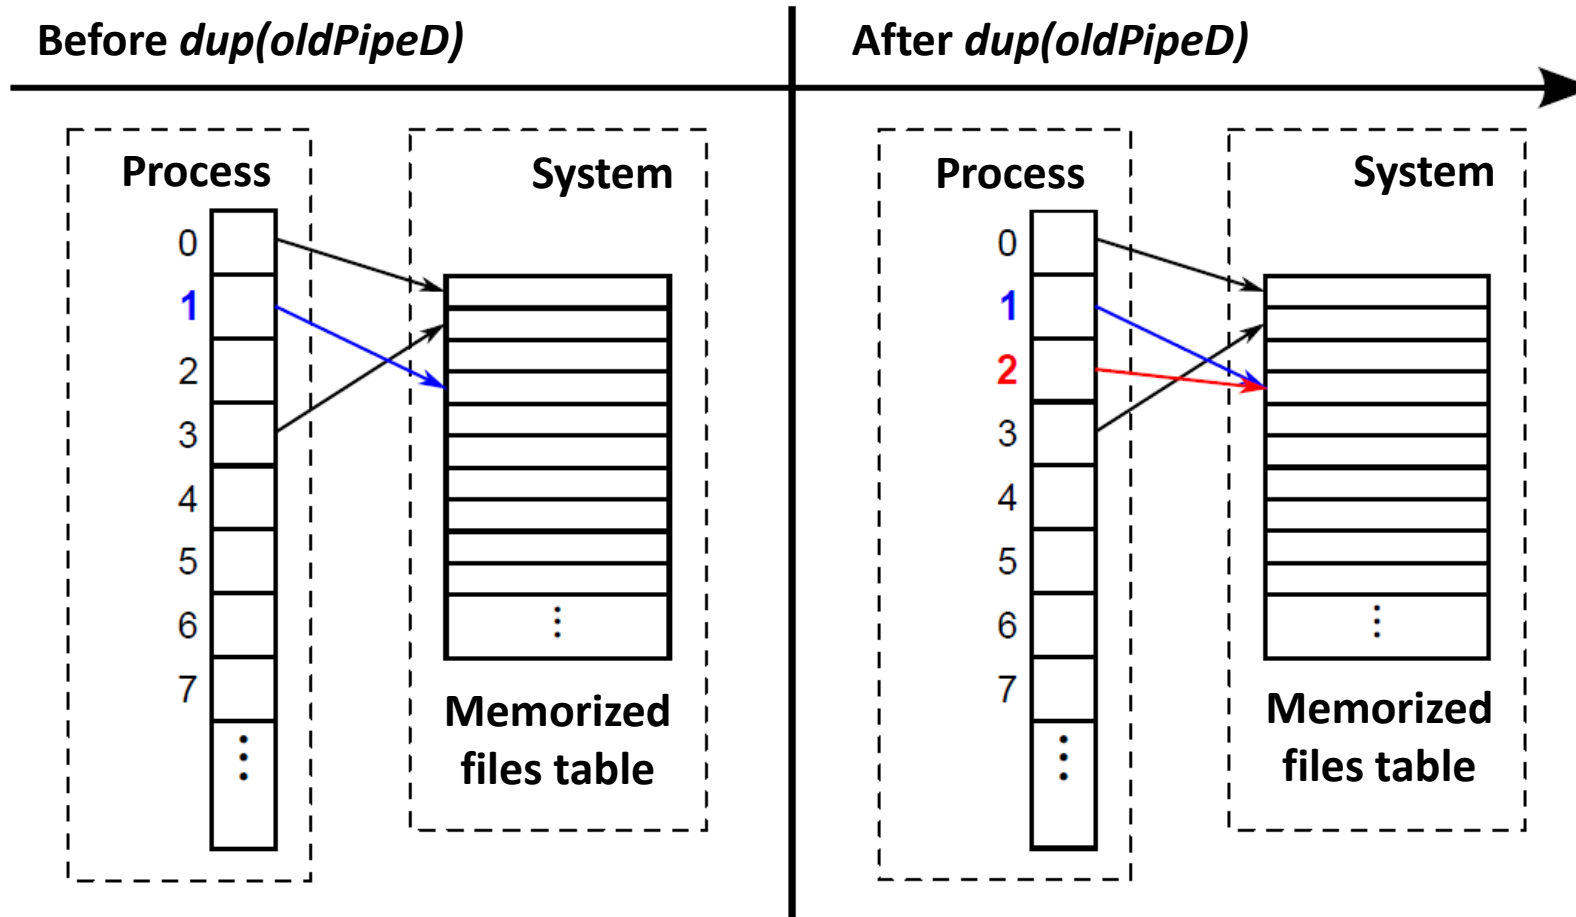
Named Pipes (3/3)

- Opening a named pipe in **blocking mode**:
 - *open()* is **reading** blocking (managing **writing**) => a process wait for another process to open the pipe in **reading** (managing **reading**)
 - The blocking opening part end in synchronous manner for the 2 processes
 - There is an automatic synchronisation of the processes opening a named pipe in blocking mode
- Opening a named pipe in **non-blocking mode**:
 - Opening in *read(O_RDONLY)* always succeed
 - Opening in *write(O_WRONLY)* only works if a process has already open the named pipe in writing
 - Writing into a pipe without reader trigger a *SIGPIPE* signal (pipe destroyed)
 - All the following read/write operations are in non-blocking mode

dup(): duplication OF a descriptor

- `#include <unistd.h>`
- `#int dup(int oldPipeD);`
- Create a copy of the file descriptor `oldPipeD`
 - Use the smallest available file descriptor
 - Return the file descriptor of the copy
- `oldPipeD` and the returned file descriptor share:
 - The position pointers of the file(*/seek*)
 - The locks, the flags (read/write, EOF, etc.) except the "close-on-exec" flag

Using dup()

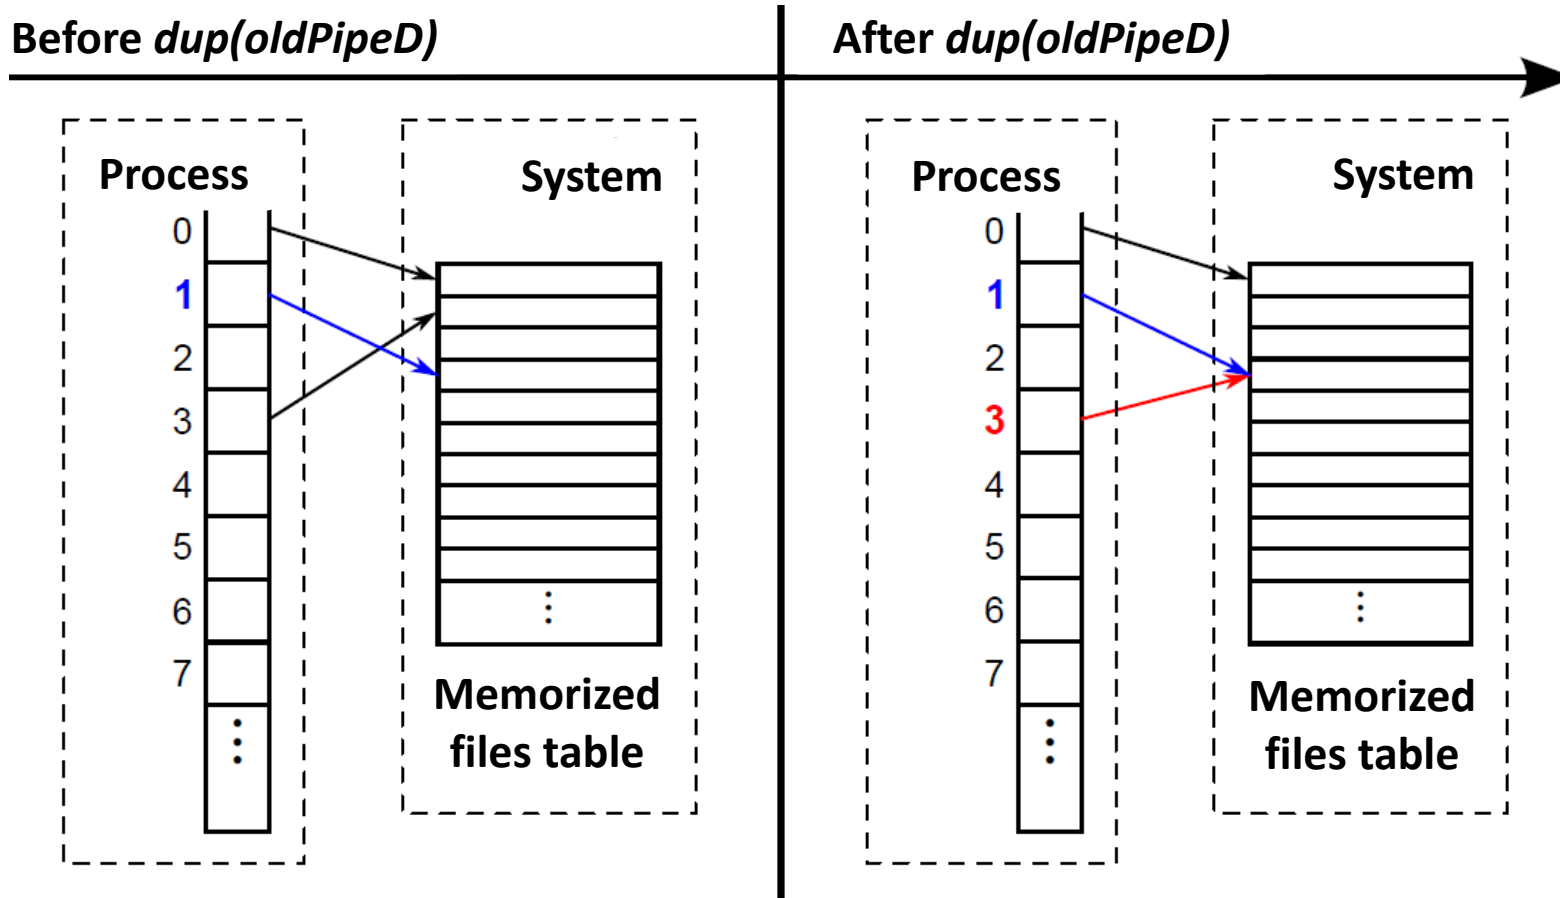


`dup (1) => 2`

dup2(): duplication ON a descriptor

- `#include <unistd.h>`
- `#int dup2(int oldPipeD, int newPipeD);`
- Works like *dup()*, apart from
 - Creating a copy of oldPipeD
 - Stops here if oldPipeD is not a valid descriptor
 - Closing newPipeD if he was opened
 - Overwriting the newPipeD with a copy of oldPipeD

Using dup2()



dup2 (1, 3) => 3

Summary: descriptor duplication (1/2)

- Using *dup()* and *dup2()*, a process can manage his file descriptors as he sees fit (for his needs)
- For example, he can associate the ending of a tube to one of his standard I/O

```
dup2(pipe[1], 1);    // Standard Output redirection  
close(pipe[1]);      // to the pipe
```

```
close(1);            // Same with dup() (descriptor 0)  
dup(pipe[1]);         // can't be free)
```

Summary: descriptor duplication (2/2)

```
saveStdout = dup(1);      // Temporary redirection of the Standard  
                           // Output to the pipe
```

```
dup2(pipe[1], 1);
```

```
close(pipe[1]);
```

```
/* ... */
```

```
dup2(saveStdout, 1);      // Standard Output Restauration
```

Exercise

- Considering a program similar to "amarok" or "rythmbox" (jukebox).
- We don't want to create ours, just to control it. Those programs creates a *named pipe* to control the process remotely. It enables us:
 - To write in the *"/tmp/commande"* pipe to send a command to the jukebox (change volume, next song, etc.)
 - To read in the *"/tmp/info"* pipe to find the current listened song trackname
- Create the following programs
 - **"currentSong"**: write on Standard Output the current listened song trackname
 - **"remoteControl"**: read a command (made of 2 integers) from Standard Output and transfer it to the jukebox